

Ridiculously Easy Centralized

Application Logging & Monitoring

Marco Pas

@marcopas

Goal

Learn how to gather logging & monitoring information from distributed systems.

House Rules



**Has the
world gone
mad?**

Or is it me?

**YOU
MAD
BRO?**





Did they come up with 911 as the police number after 9/11? #wondering

Reply Retweeted Favorited Buffer

 Follow



< Tweet 🔍 ✍️



 +

I'm sick of the U.S Government. I'm moving to California. Fuck the U.S.

10/1/13, 12:44 AM



John 

 +

Gonna stop smoking for a week.
Feel bad for my liver :S

2012-08-27 6:25 AM











1st law of distributed computing

“Do not distribute
until you really need it”

Let's make the world easier by using... Distributed Computing



Monolith



Microservices

It keeps the code cleaner

“You don’t need to introduce a network boundary
as an excuse to write better code”

It's easy to write things that
only have one purpose

“Distributed Transactions are never easy”

They're faster than monoliths

“You could gain a lot of performance in a monolith by simply applying a little extra discipline”

It's easy for engineers to not all work in the same codebase

“A bunch of engineers working in isolated codebases leads to ‘Not my problem’ syndrome”

It's the simplest way to
handle autoscaling, plus
Docker is in here somewhere

“You can scale a microservice outward just as easily
as you can scale a monolith”

DISTRIBUTED SYSTEMS




HOW HARD CAN IT BE?

Logging



- Providing useful information, seems hard!
- Common Log Formats
 - W3C, Common Log Format, Combined Log Format
 - used for:
 - Proxy & Web Servers
- Agree upon Application Log Formats
 - Do not forget -> Log levels!
- Data security
 - Do not log passwords or privacy related data



Some seriously useful log message :)

- “No need to log, we know what is happening”
- “Something happened not sure what”
- “Empty log message”
- “Lots of sh*t happing”
- “It works b****”
- “How did we end up here?”
- “Okay i am getting tired of this error message”
- “Does this work?”
- “We hit a bug, still figuring out what”
- “Call 911 we have a problem”



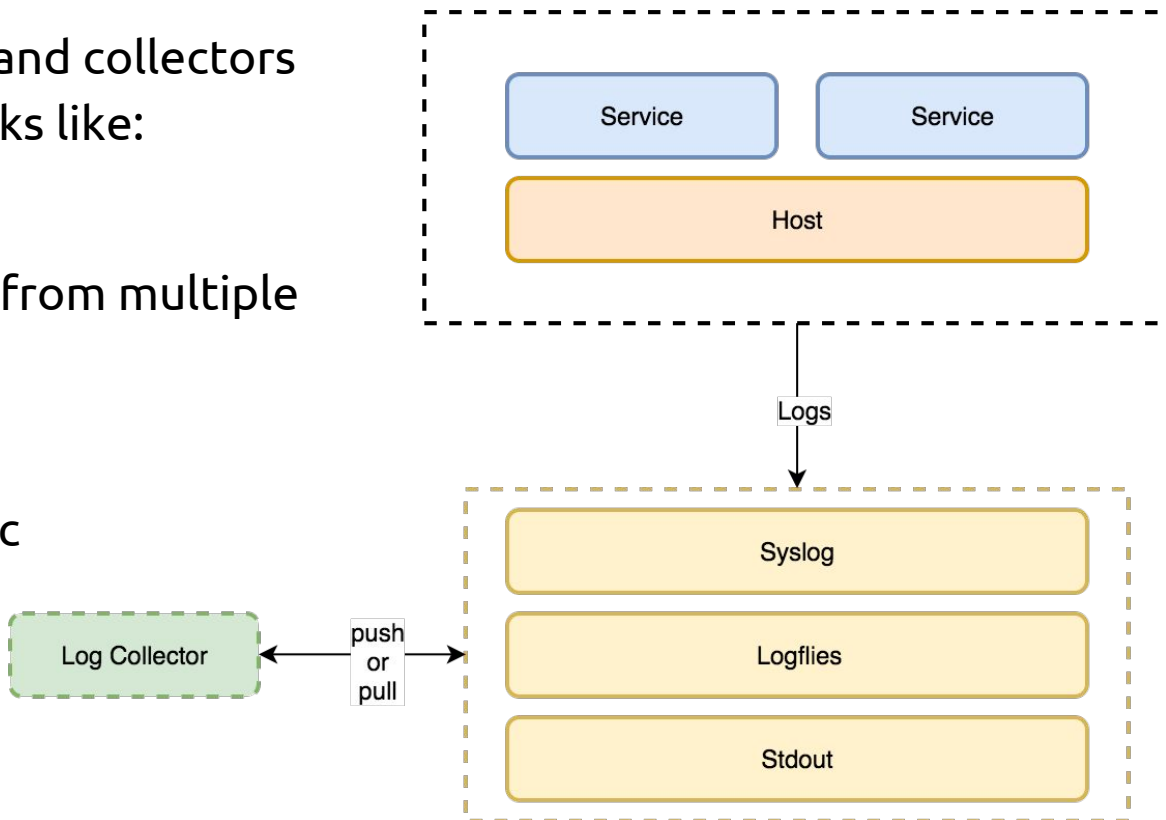
- Syslog / Syslog-ng
- Files -> multiple places (/var/log)
 - Near real time replication to remote destinations
- Stdout
 - Normally goes to /dev/null



In container based environments logging to “Stdout” has the preference



- Specialized transporters and collectors available using frameworks like:
 - Logstash, Flume, Fluentd
- Accumulate data coming from multiple hosts / services
 - Multiple input sources
- Optimized network traffic
 - Pull / Push






- Where should it be stored?

- Short vs Long term
- Associated costs
- Speed of data ingestion & retrieval
- Data access policies (who needs access)

- Example storage options:

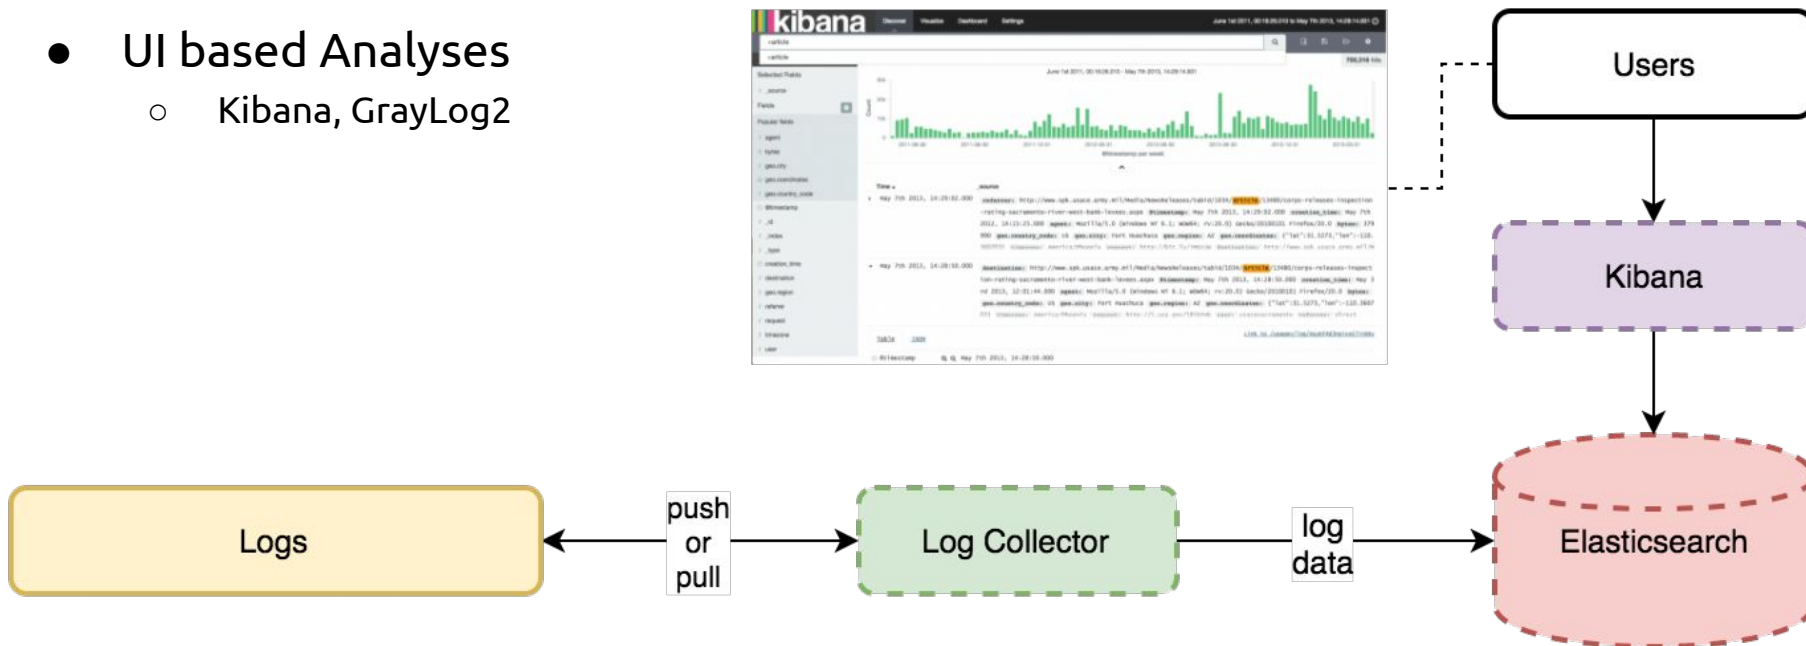
- S3, Glacier, Tape backup
- HDFS, Cassandra, MongoDB or ElasticSearch





- Batch processing of log data
 - HDFS, Hive, PIG → MapReduce Jobs

- UI based Analyses
 - Kibana, GrayLog2





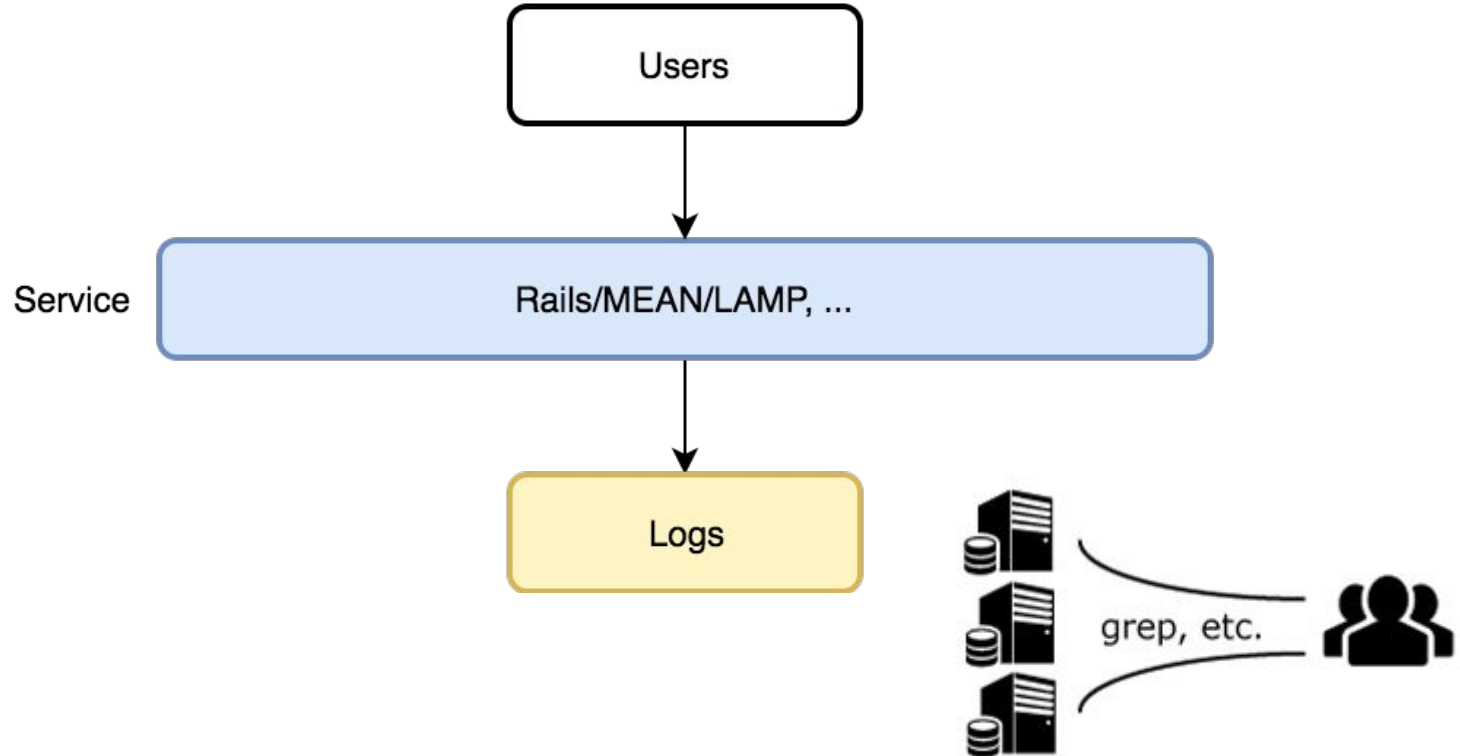
- Based on patterns or “calculated” metrics → send out events
 - Trigger alert and send notifications
- Logging != Monitoring
 - Logging -> recording to diagnose a system

```
127.0.0.1 - frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0" 200 2326
```

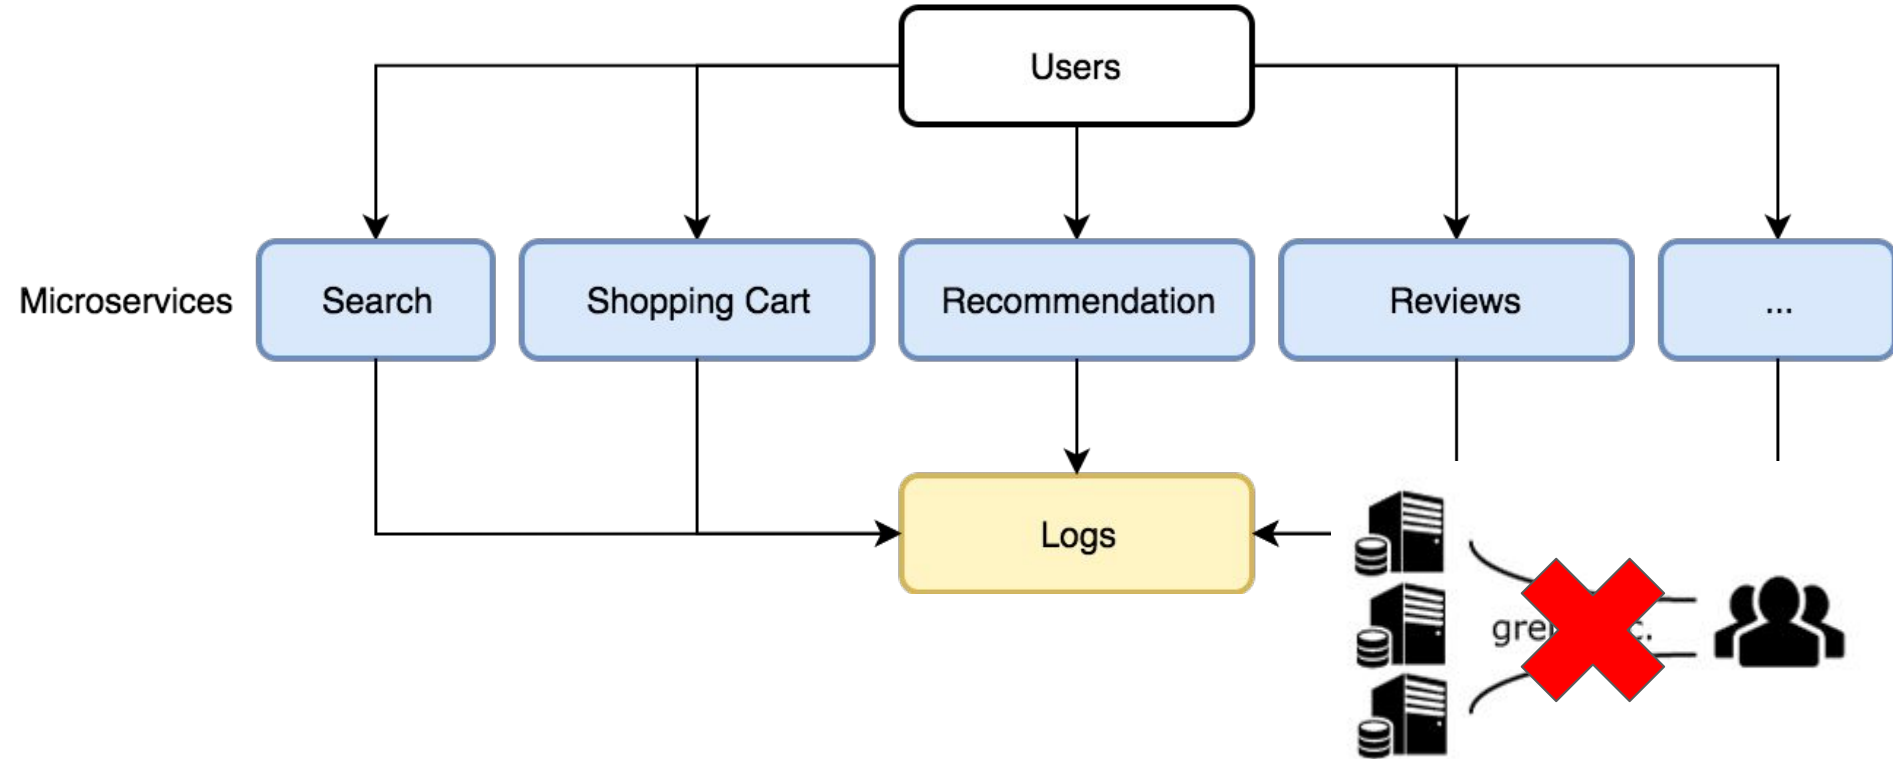
- Monitoring -> observation, checking and recording

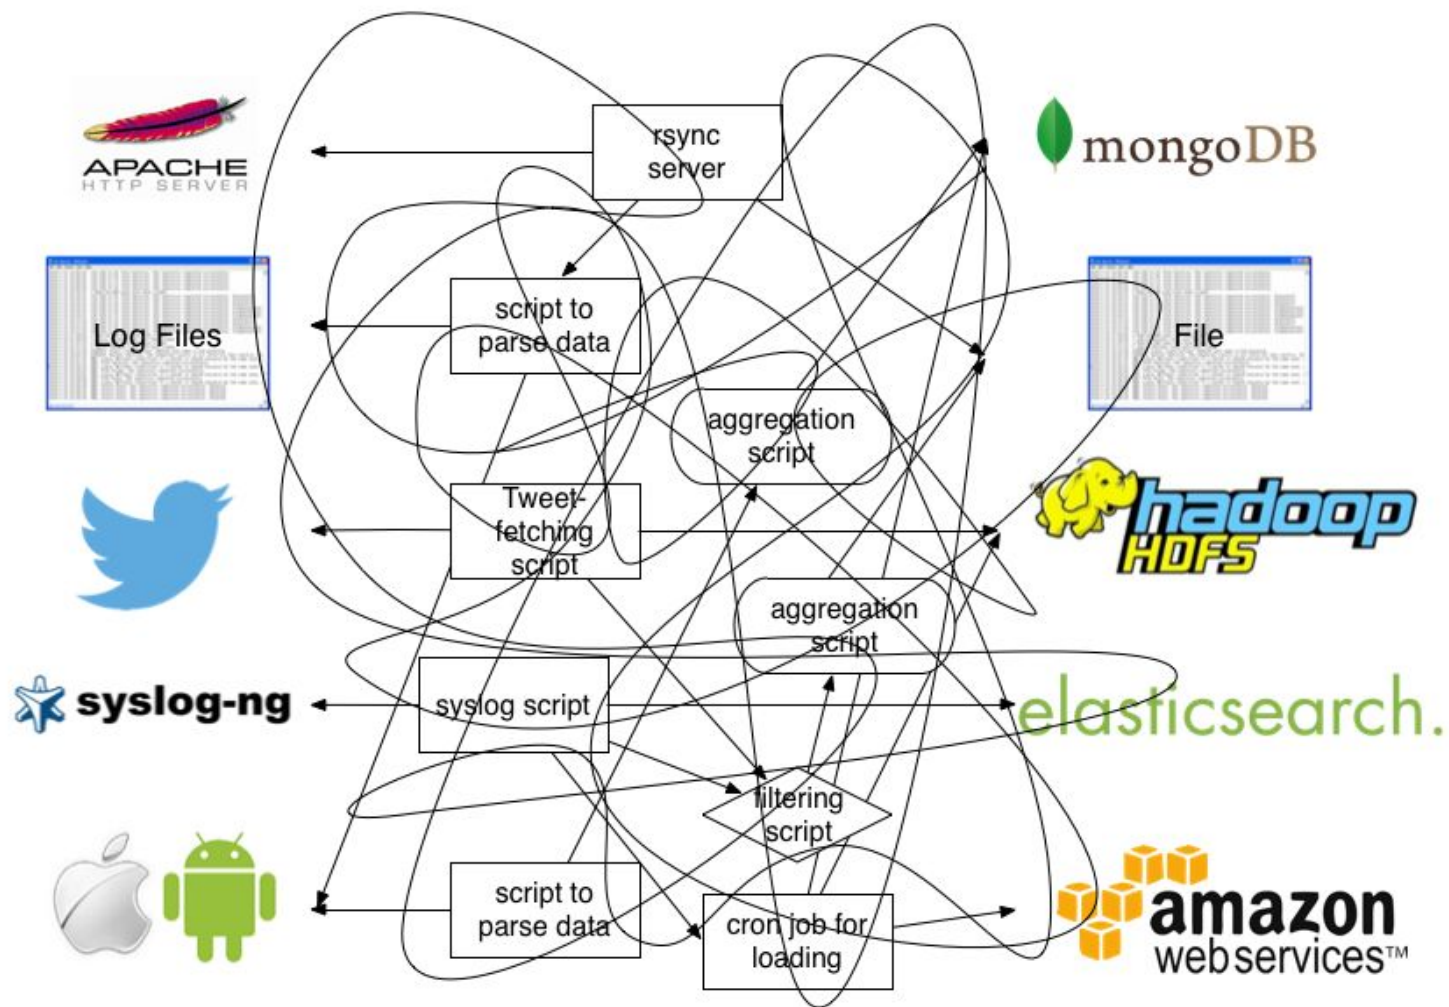
```
http_requests_total{method="post",code="200"} 1027 1395066363000
```

Logging




Distributed Logging





Need for a Unified Logging Layer



Access logs

Apache

App logs

Frontend
Backend

System logs

syslogd

Databases

Alerting

Nagios

Analysis

MongoDB
MySQL
Hadoop

Archiving

Amazon S3




fluentd



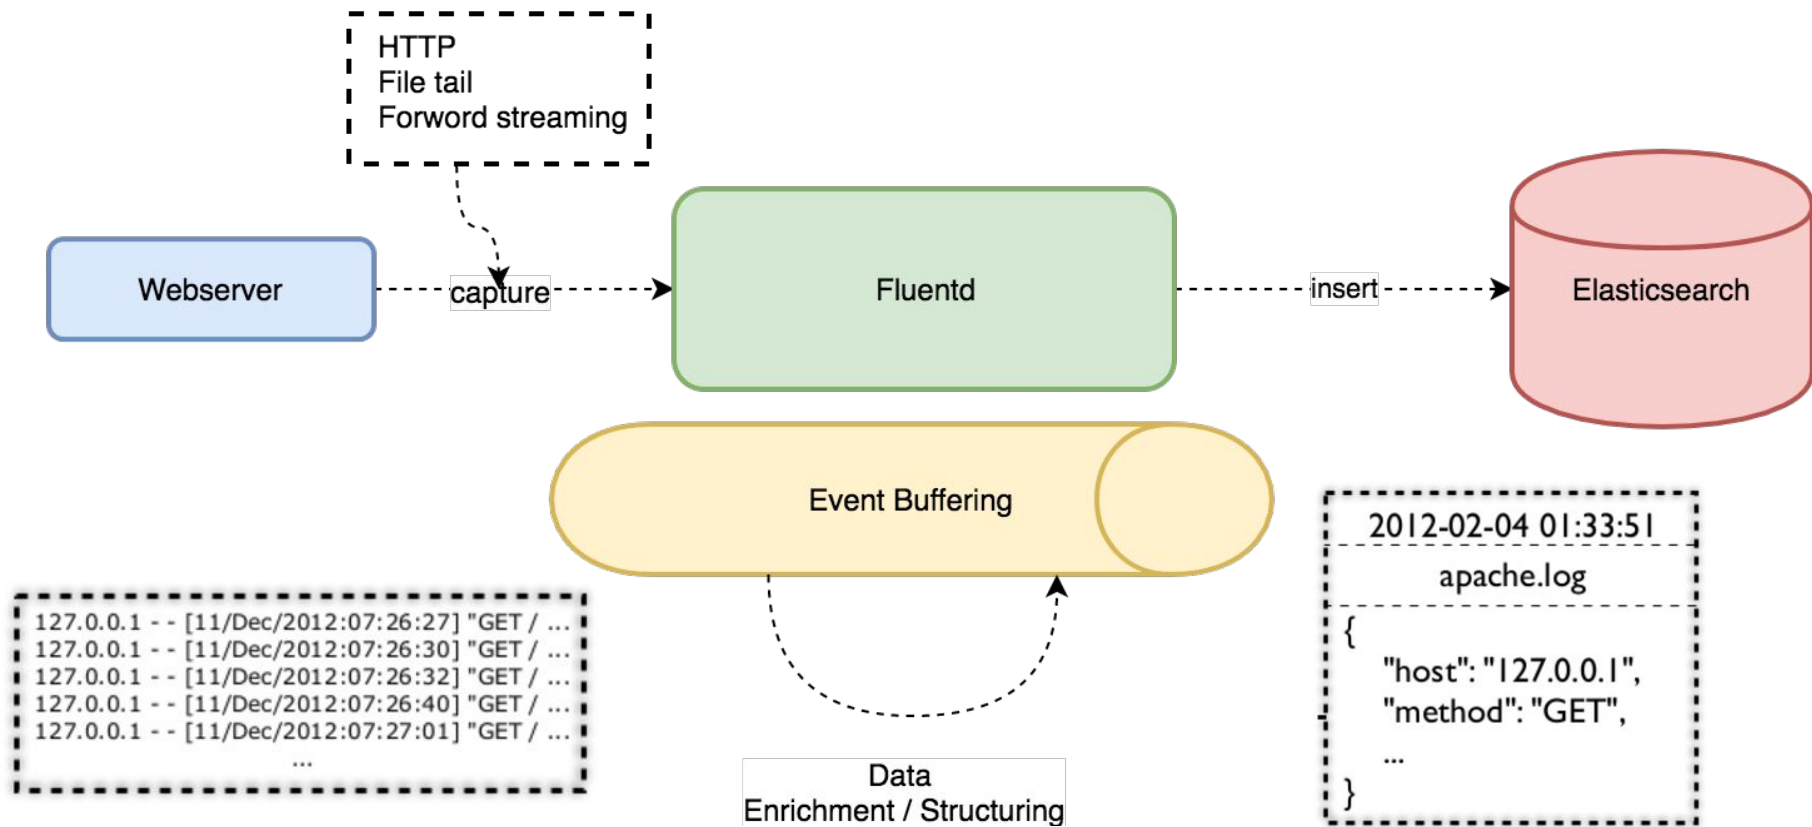
filter / buffer / routing

Fluentd

- Open source log collector written in Ruby
- Reliable, scalable and easy to extend
 - Pluggable architecture
 - Rubygem ecosystem for plugins
- Reliable log forwarding

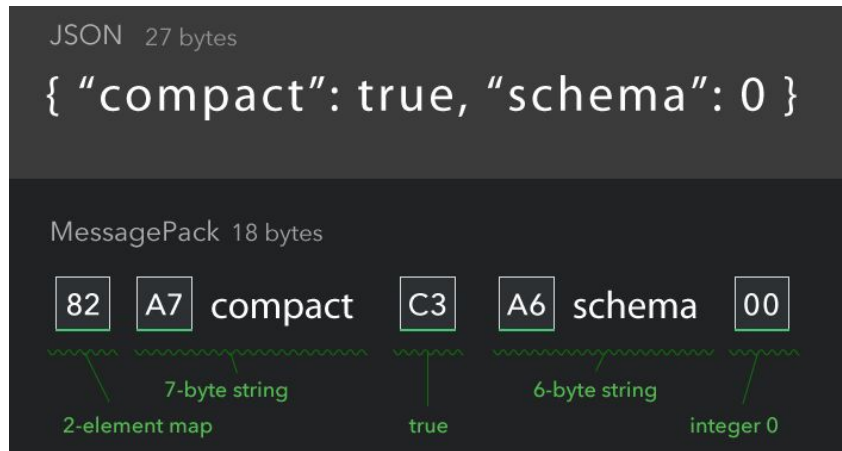


Example



Event structure

- Tag
 - Where an event comes from, used for message routing
- Time
 - When an event happens, Epoch time
 - Parsed time coming from the datasource
- Record
 - Actual log content being a JSON object
 - Internally MessagePack



Event example

```
192.168.0.1 - - [28/Feb/2013:12:00:00 +0900] "GET / HTTP/1.1" 200 777
```



```
tag:: apache.access # set by configuration  
time: 1362020400    # 28/Feb/2013:12:00:00 +0900  
record: {"user":"-","method":"GET","code":200,"size":777,"host":"192.168.0.1","path":"/"}
```


Configuration

- Driven by a simple text based configuration file
 - fluent.conf

```
<source></source>
```

→ Tell where the data comes from (input)

```
<match></match>
```


→ Tell fluentd what to do (output)

```
<filter></filter>
```

→ Event processing pipeline

```
<label></label>
```

→ Groups filter and output for internal routing



source -> filter 1 -> ... -> filter N -> output

```
# receive events via HTTP
```

```
<source>
```

```
  @type http
```

```
  port 9880
```

```
</source>
```

```
# read logs from a file
```

```
<source>
```

```
  @type tail
```

```
  path /var/log/httpd.log
```

```
  format apache
```

```
  tag apache.access
```

```
</source>
```

```
# add a field to an event
```

```
<filter myapp.access>
```

```
  @type record_transformer
```

```
  <record>
```

```
    host_param "#{Socket.gethostname}"
```

```
  </record>
```

```
</filter>
```

```
# save access logs to MongoDB
```

```
<match apache.access>
```

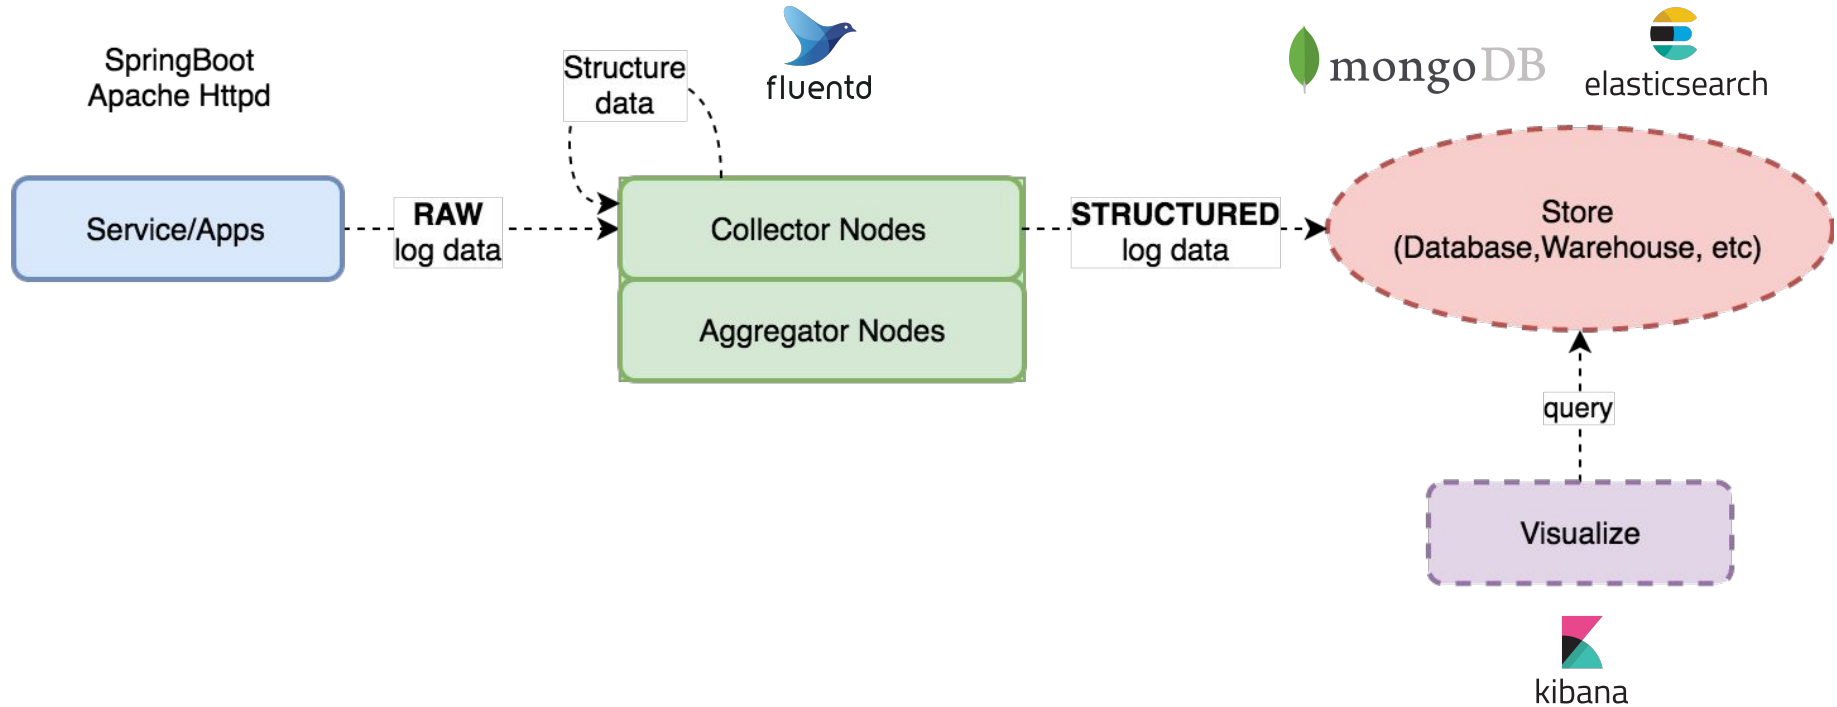
```
  @type mongo
```

```
  database apache
```

```
  collection log
```

```
</match>
```

Demo: Capture Grails/Spring Boot Logs



Code Demo

“Capture Grails/Spring Boot Logs”


Monitoring

Our scary movie “The Happy Developer”

- Let's push out features
- I can demo so it works :)
- It works with 1 user, so it will work with multiple
- Don't worry about performance we will just scale using multiple machines/processes
- Logging is into place



Disaster Strikes



Logging != Monitoring

Logging

“recording to diagnose a system”

```
127.0.0.1 - frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0" 200 2326
```

Monitoring

“observation, checking and recording”

```
http_requests_total{method="post",code="200"} 1027 1395066363000
```


Vital Signs

14:21

Unit:

HR
bpm

200
90

154
(bpm)

NIBP
mmHg

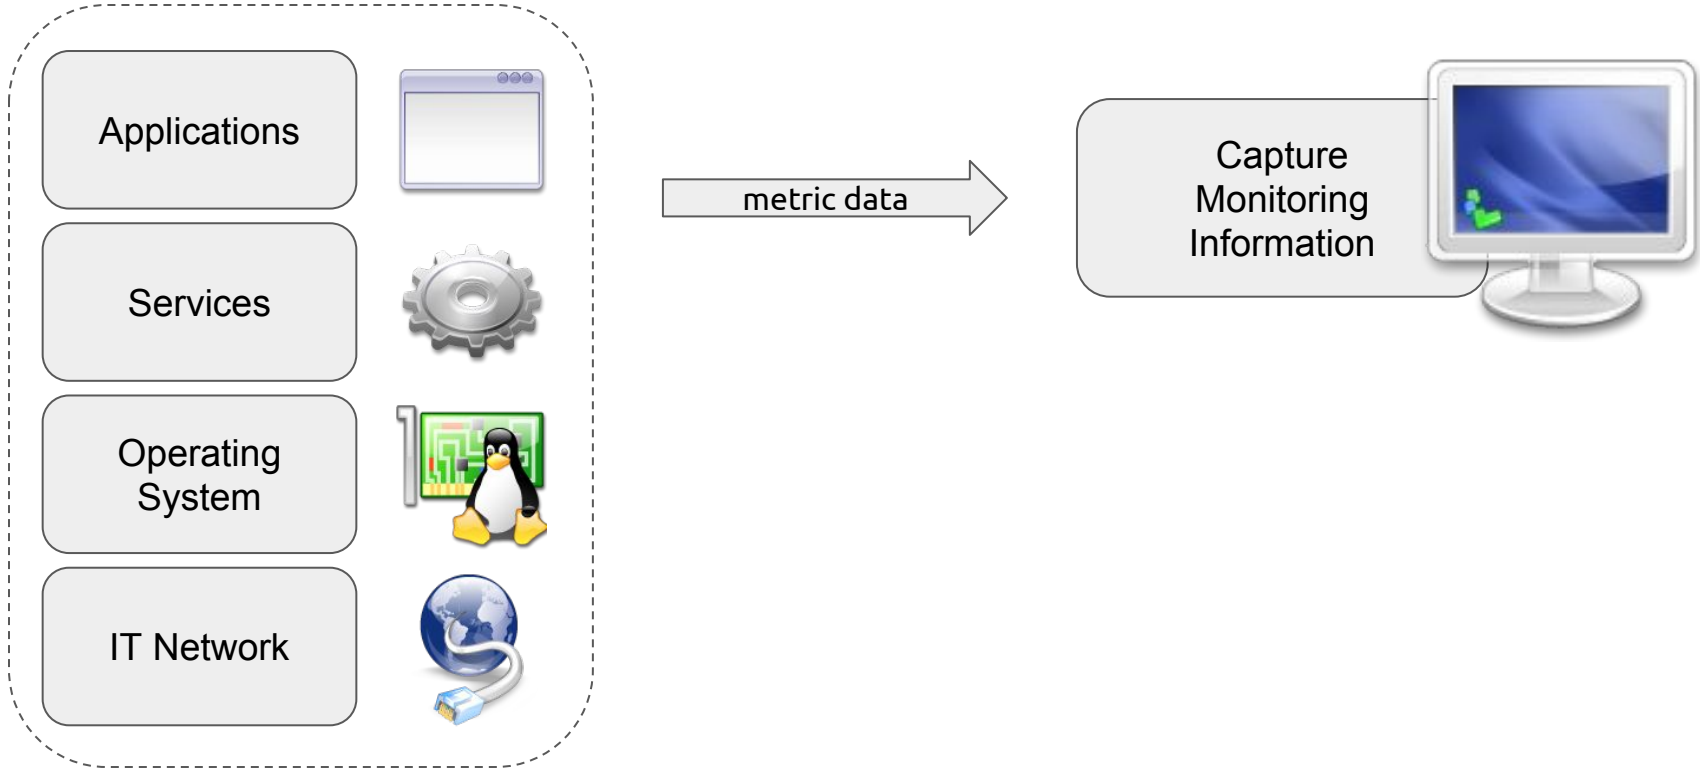
cuff: 0

Why Monitoring?



- Know when things go wrong
 - Detection & Alerting
- Be able to debug and gain insight
- Detect changes over time and drive technical/business decisions
- Feed into other systems/processes (e.g. security, automation)



What to monitor?



Houston we have Storage problem!





How to store the mass amount of metrics and also making them easy to query?

Time Series - Database

- Time series data is a sequence of data points collected at regular intervals over a period of time. (metrics)

- Examples:

- Device data
- Weather data
- Stock prices
- Tide measurements
- Solar flare tracking



- The data requires aggregation and analysis

Time Series - Data format

metric name and a set of key-value pairs, also known as labels



















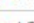






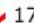




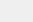
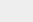


`<metric name>{<label name>=<label value>, ...} value [timestamp]`



`http_requests_total{method="post",code="200"} 1027 1395066363000`

23 systems in ranking, May 2018

Rank			DBMS	Database Model	Score		
May 2018	Apr 2018	May 2017			May 2018	Apr 2018	May 2017
1.	1.	1.	InfluxDB 	Time Series DBMS	11.00	+0.24	+3.05
2.	2.	 5.	Kdb+ 	Multi-model 	3.07	-0.01	+1.50
3.	3.	 2.	RRDtool	Time Series DBMS	2.68	-0.07	-0.34
4.	4.	 3.	Graphite	Time Series DBMS	2.26	+0.07	+0.25
5.	5.	 4.	OpenTSDB	Time Series DBMS	1.62	-0.08	-0.05
6.	 7.	 8.	Prometheus	Time Series DBMS	1.12	+0.07	+0.64
7.	 6.	 6.	Druid	Time Series DBMS	1.01	-0.05	+0.07
8.	8.	 7.	KairosDB	Time Series DBMS	0.43	-0.01	-0.08
9.	9.	9.	eXtremeDB 	Multi-model 	0.31	-0.01	-0.02
10.	10.	 11.	Riak TS	Time Series DBMS	0.27	-0.00	+0.06
11.	11.	 19.	Hawkular Metrics	Time Series DBMS	0.11	+0.00	+0.11
12.	12.	 18.	Blueflood	Time Series DBMS	0.10	+0.00	+0.07
13.	 15.	 10.	Axibase	Time Series DBMS	0.05	+0.02	-0.16
14.	 18.	 12.	Warp 10	Time Series DBMS	0.04	+0.01	-0.15
15.	 16.		TimescaleDB	Time Series DBMS	0.04	+0.01	
16.	 17.	 13.	TempoIQ	Time Series DBMS	0.02	-0.00	-0.13
17.	 13.	 15.	Machbase 	Time Series DBMS	0.01	-0.07	-0.04
18.	 19.	 17.	Heroic	Time Series DBMS	0.00	±0.00	-0.03
18.	 19.		IRONdb	Time Series DBMS	0.00	±0.00	
18.	 19.	 19.	Newts	Time Series DBMS	0.00	±0.00	±0.00



<http://db-engines.com/en/ranking/time+series+dbms>

Prometheus Overview

Prometheus


Prometheus is an open-source systems monitoring and alerting toolkit originally built at SoundCloud. It is now a standalone open source project and maintained independently of any company.

<https://prometheus.io>

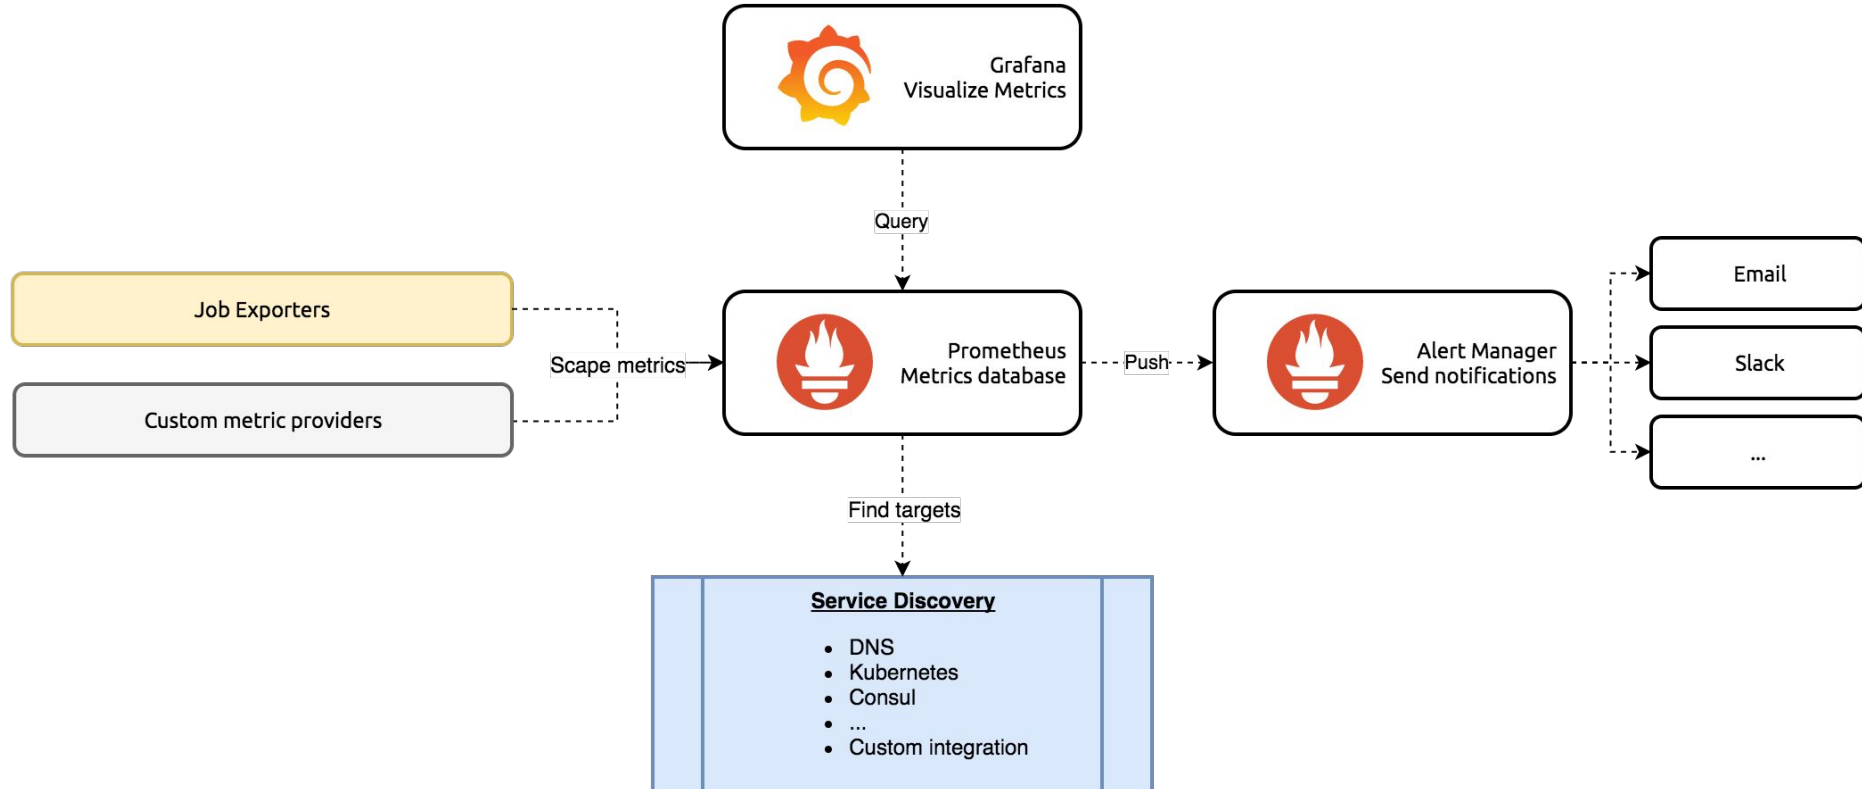


Prometheus Components

- The main [Prometheus server](#) which scrapes and stores time series data
- [Client libraries](#) for instrumenting application code
- A [push gateway](#) for supporting short-lived jobs
- Special-purpose [exporters](#) (for HAProxy, StatsD, Graphite, etc.)
- An [alertmanager](#)
- Various support tools



Prometheus Overview



List of Job Exporters

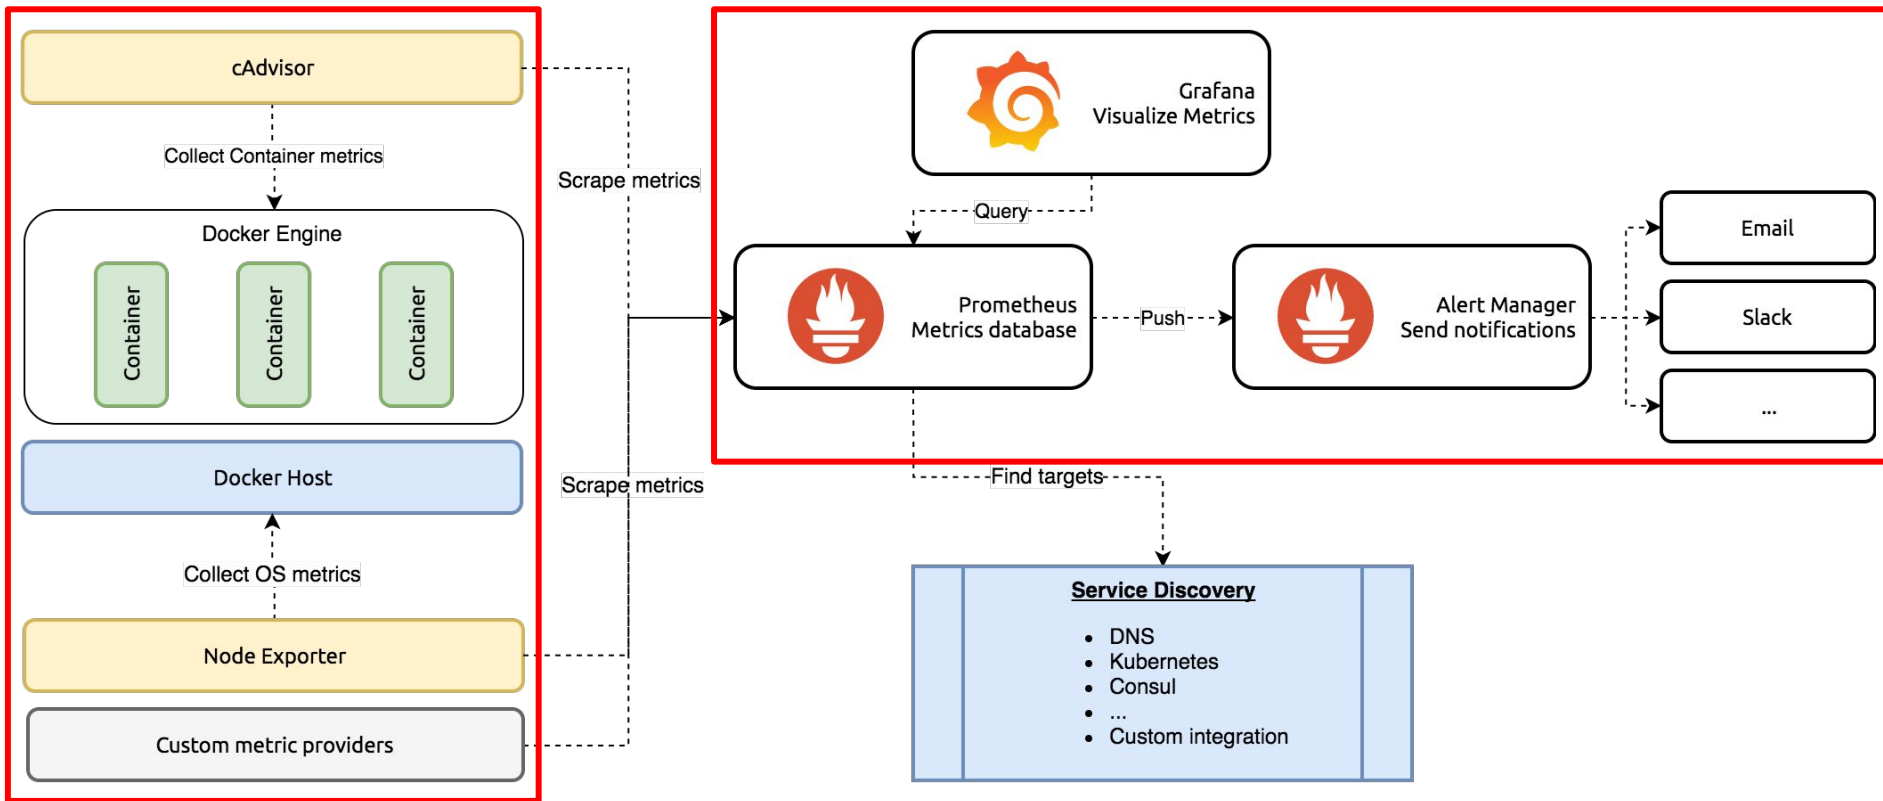
- Prometheus managed:

- JMX
- Node
- Graphite
- Blackbox
- SNMP
- HAProxy
- Consul
- Memcached
- AWS Cloudwatch
- InfluxDB
- StatsD
- ...

- Custom ones:

- Database
- Hardware related
- Messaging systems
- Storage
- HTTP
- APIs
- Logging
- ...

Demo: Application Monitoring



Code Demo

“Prometheus monitoring
including alerting”

That's a wrap!

Question?